


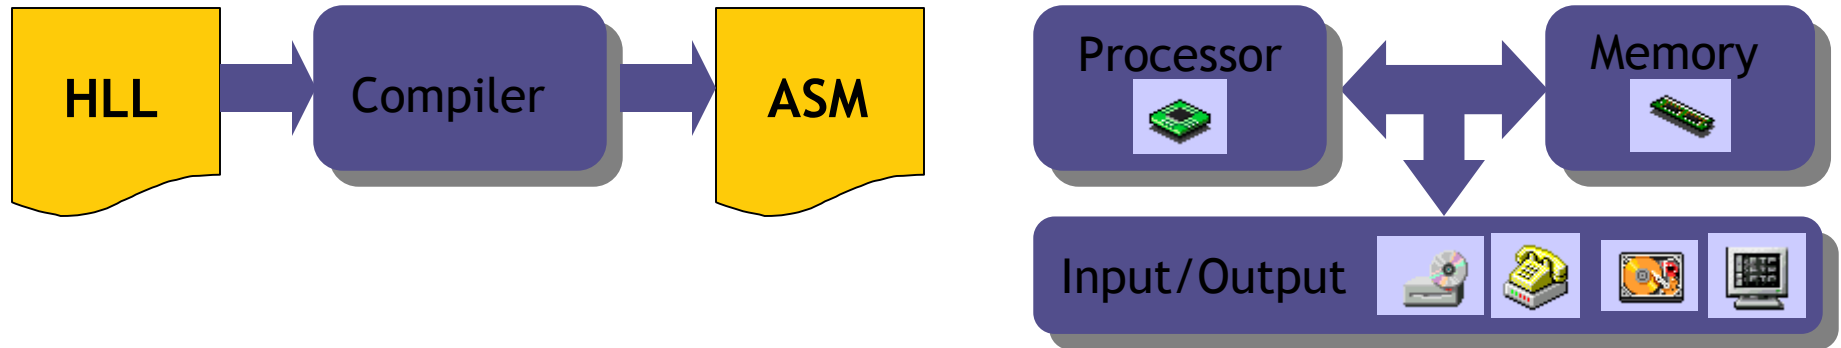
CS232: Computer Architecture II

Fall 2006




What is computer architecture about?

- Computer architecture is about building and analyzing computer systems.




- In CS232, we will take a tour of the whole machine.
- Specifically, we'll...

Study Instruction Set Architectures



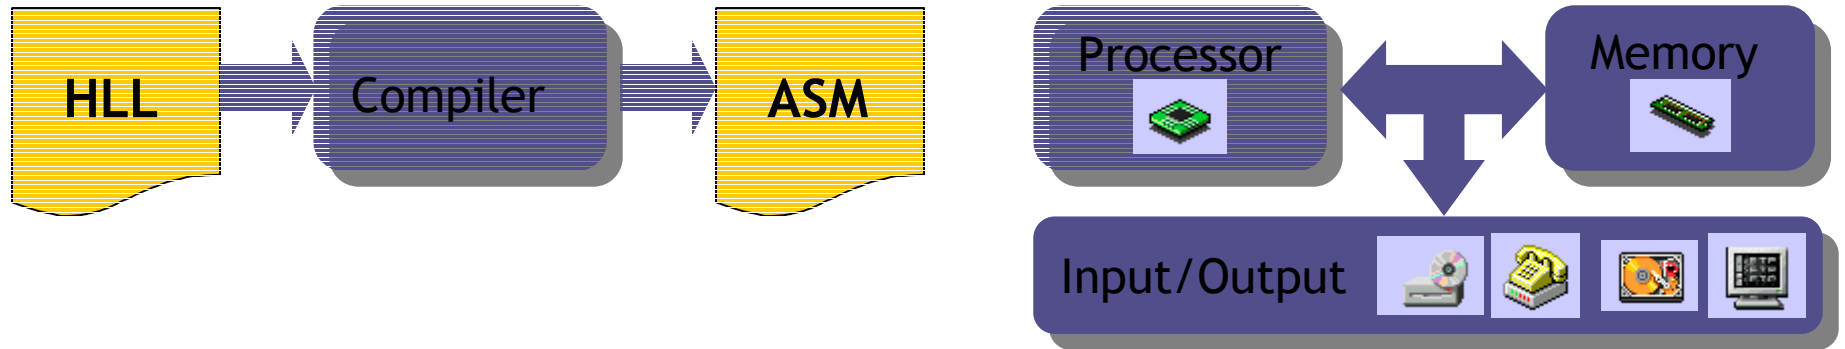
- The Instruction Set Architecture (ISA) is the bridge between the hardware and the software.
- We'll learn the MIPS ISA in detail
- We'll learn how HLL program constructs are represented to the machine
- We won't learn how compilers work, but we'll learn what they do

Learn about Modern Processor Organization



- What makes a processor fast?
- The key technique we'll focus on is: Pipelining
 - Pipelining allows processors to work on multiple instructions at the same time.

Learn about Memory and I/O systems




- We'll learn how virtual memory makes programming easy
- We'll learn how caches make memory fast
- We'll learn about buses and disks

Why should you care?

- It is **interesting**.
 - How do you make a processor that runs at 4Ghz?
- It will help you be a **better programmer**.
 - Understanding how your program is translated to assembly code lets you reason about correctness and performance.
 - Demystify the seemingly arbitrary (*e.g.*, bus errors, segmentation faults)
- Many **cool jobs** require an understanding of computer architecture.
 - The cutting edge is often pushing computers to their limits.
 - Supercomputing, games, portable devices, etc.
- Computer architecture illustrates many **fundamental ideas** in computer science
 - Abstraction, caching, and indirection are everywhere in CS

CS231 vs. CS232

- This class expands upon the computer architecture material from the last few weeks of CS231, and we rely on many other ideas from CS231.
 - Understanding binary, hexadecimal and two's-complement numbers is still important.
 - Devices like multiplexers, registers and ALUs appear frequently. You should know what they do, but not necessarily how they work.
 - Finite state machines and sequential circuits will appear again.
- We do *not* spend much time with logic design topics like Karnaugh maps, Boolean algebra, latches and flip-flops.



Who we are

- *Lecturer:*

Viraj Kumar

- I'm a graduate student interested in teaching

- *Teaching Assistants:*


Abhilasha Chaudhary

Vijay Nori

Peter Young

- *Office:* 0212 Siebel


Instruction set architectures



- In CS232, we'll talk about several important issues that we didn't see in the simple processor from CS231.
 - The instruction set in CS231 lacked many features, such as support for function calls. We'll work with a larger, more realistic processor.
 - We'll also see more ways in which the instruction set architecture affects the hardware design.

MIPS

- In this class, we'll use the MIPS instruction set architecture (ISA) to illustrate concepts in assembly language and machine organization
 - Of course, the concepts are not MIPS-specific
 - MIPS is just convenient because it is real, yet simple (unlike x86)
- The MIPS ISA is still used in many places today. Primarily in embedded systems, like:
 - Various routers from [Cisco](#)
 - Game machines like the [Nintendo 64](#) and [Sony Playstation 2](#)



What you will need to learn this month

- You must become “fluent” in MIPS assembly:
 - Translate from C to MIPS and MIPS to C
- Example problem a previous last mid-term 1:

Question 3: Write a recursive function (30 points)


Here is a function `pow` that takes two arguments (`n` and `m`, both 32-bit numbers) and returns n^m (i.e., `n` raised to the m^{th} power).

```
int
pow(int n, int m) {
    if (m == 1)
        return n;
    return n * pow(n, m-1);
}
```

Translate this into a MIPS assembly language function.

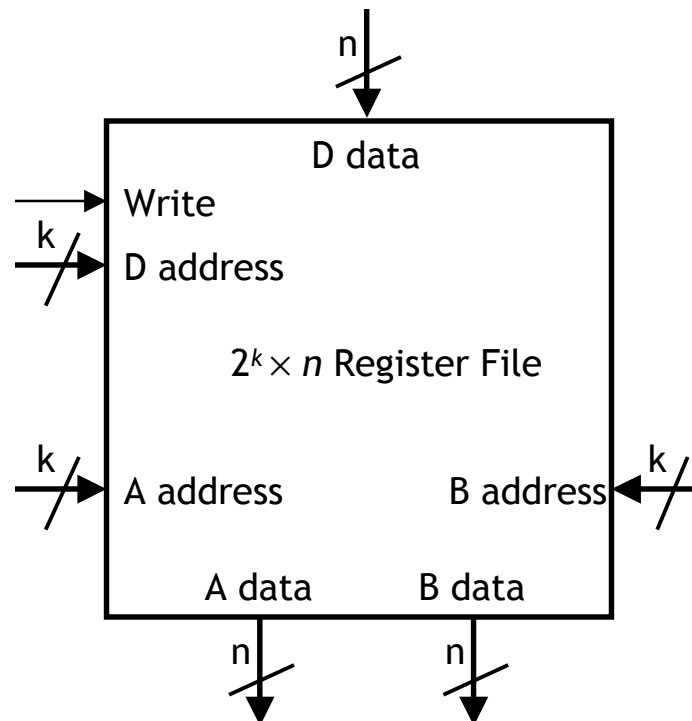
MIPS: register-to-register, three address

- MIPS is a **register-to-register**, or **load/store**, architecture.
 - The destination and sources must all be registers.
 - Special instructions, which we'll see later today, are needed to access main memory.
- MIPS uses **three-address** instructions for data manipulation.
 - Each ALU instruction contains a **destination** and two **sources**.
 - For example, an addition instruction ($a = b + c$) has the form:



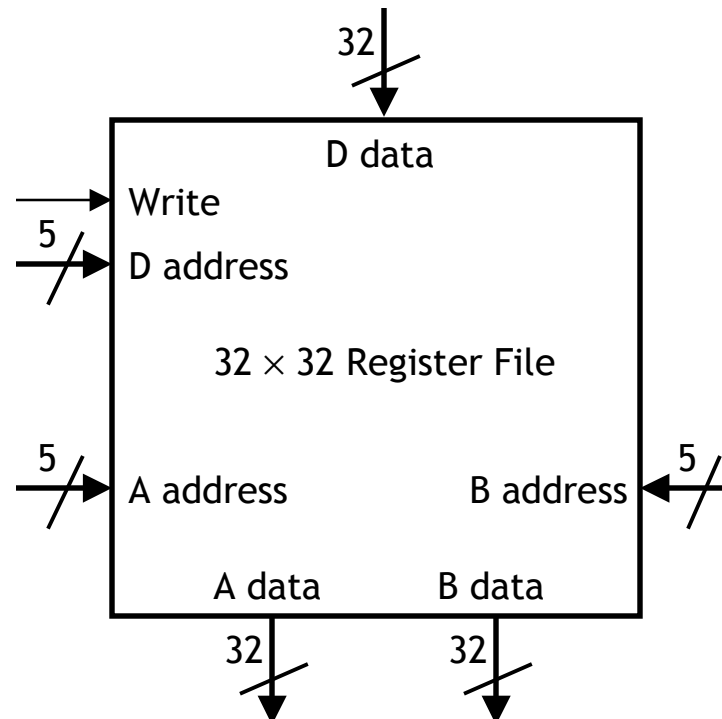
Register file review

- Here is a block symbol for a general $2^k \times n$ register file.
 - If **Write** = 1, then **D data** is stored into **D address**.
 - You can read from two registers at once, by supplying the **A address** and **B address** inputs. The outputs appear as **A data** and **B data**.
- Registers are clocked, sequential devices.
 - We can read from the register file at any time.
 - Data is written only on the positive edge of the clock.



MIPS register file

- MIPS processors have 32 registers, each of which holds a 32-bit value.
 - Register addresses are 5 bits long.
 - The data inputs and outputs are 32-bits wide.
- More registers might seem better, but there is a limit to the goodness.
 - It's more expensive, because of both the registers themselves as well as the decoders and muxes needed to select individual registers.
 - Instruction lengths may be affected, as we'll see in the future.



MIPS register names

- MIPS register names begin with a **\$**. There are two naming conventions:
 - By number:

`$0 $1 $2 ... $31`

- By (mostly) two-character names, such as:

`$a0-$a3 $s0-$s7 $t0-$t9 $sp $ra`

- Not all of the registers are equivalent:
 - E.g., register `$0` or `$zero` always contains the value 0
 - (go ahead, try to change it)
- Other registers have special uses, by convention:
 - E.g., register `$sp` is used to hold the “stack pointer”
- You have to be a little careful in picking registers for your programs
 - for now, stick to the registers `$t0-$t9`

Basic arithmetic and logic operations

- The basic integer arithmetic operations include the following:


add sub mul div

- And here are a few logical operations:

and or xor nor

- Remember that these all require three register operands; for example:

```
add  $t0, $t1, $t2      # $t0 = $t1 + $t2  
mul  $s1, $s1, $a0      # $s1 = $s1 x $a0
```



Larger expressions

- More complex arithmetic expressions may require multiple operations at the instruction set level.

$$t0 = (t1 + t2) \times (t3 - t4)$$

```
add $t0, $t1, $t2    # $t0 contains $t1 + $t2
sub $s0, $t3, $t4    # Temporary value $s0 = $t3 - $t4
mul $t0, $t0, $s0    # $t0 contains the final product
```

- Temporary registers may be necessary, since each MIPS instructions can access only two source registers and one destination.
 - In this example, we could re-use \$t3 instead of introducing \$s0.
 - But be careful not to modify registers that are needed again later.

Immediate operands

- The ALU instructions we've seen so far expect register operands. How do you get data into registers in the first place?
 - Some MIPS instructions allow you to specify a signed constant, or “immediate” value, for the second source instead of a register. For example, here is the immediate add instruction, **addi**:

```
addi $t0, $t1, 4      # $t0 = $t1 + 4
```

- Immediate operands can be used in conjunction with the **\$zero** register to write constants into registers:

```
addi $t0, $0, 4       # $t0 = 4
```

Shorthand:

```
li $t0 4
```

\$t0 = 4
(pseudo-instruction)


- MIPS is still considered a load/store architecture, because arithmetic operands cannot be from arbitrary memory locations. They must either be registers or constants that are embedded in the instruction.

We need more space!

- Registers are fast and convenient, but we have only 32 of them, and each one is just 32-bits wide.
 - That's not enough to hold data structures like large arrays.
 - We also can't access data elements that are wider than 32 bits.
- We need to add some main memory to the system!
 - RAM is cheaper and denser than registers, so we can add lots of it.
 - But memory is also significantly slower, so registers should be used whenever possible.
- In the past, using registers wisely was the programmer's job.
 - For example, C has a keyword “register” that marks commonly-used variables which should be kept in the register file if possible.
 - However, modern compilers do a pretty good job of using registers intelligently and minimizing RAM accesses.

Memory review


- Memory sizes are specified much like register files; here is a $2^k \times n$ RAM.




CS	WR	Operation
0	x	None
1	0	Read selected address
1	1	Write selected address

- A chip select input **CS** enables or “disables” the RAM.
- ADRS** specifies the memory location to access.
- WR** selects between reading from or writing to the memory.
 - To read from memory, **WR** should be set to 0. **OUT** will be the n-bit value stored at **ADRS**.
 - To write to memory, we set **WR** = 1. **DATA** is the n-bit value to store in memory.

MIPS memory



- MIPS memory is **byte-addressable**, which means that each memory address references an 8-bit quantity.
- The MIPS architecture can support up to 32 address lines.
 - This results in a $2^{32} \times 8$ RAM, which would be 4 GB of memory.
 - Not all actual MIPS machines will have this much!



Loading and storing bytes

- The MIPS instruction set includes dedicated load and store instructions for accessing memory, much like the CS231 example processor.
- The main difference is that MIPS uses **indexed addressing**.
 - The address operand specifies a signed constant and a register.
 - These values are added to generate the effective address.
- The MIPS “load byte” instruction **lb** transfers one byte of data from main memory to a register.

```
lb $t0, 20($a0)      # $t0 = Memory[$a0 + 20]
```

- The “store byte” instruction **sb** transfers the lowest byte of data from a register into main memory.

```
sb $t0, 20($a0)      # Memory[$a0 + 20] = $t0
```

Byte loads

- **Question:** if you load a byte (8 bits) into a register (32 bits), what value do those other 24 bits have?

Consider the code mem.s demonstrated in class (see website)

`$t0` is initialized to 0x0000FFFF

if the byte 0x08 is loaded into `$t0` (using `lb`), then

`$t0` = 0x00000008

however, if the byte 0x80 is loaded into `$t0`, then

`$t0` = 0xFFFFFFFF80


What is going on here?? Answer: sign-extension!

Loading and storing words

- You can also load or store 32-bit quantities—a complete **word** instead of just a byte—with the **lw** and **sw** instructions.

```
lw $t0, 20($a0)      # $t0 = Memory[$a0 + 20]
sw $t0, 20($a0)      # Memory[$a0 + 20] = $t0
```

- Most programming languages support several 32-bit data types.
 - Integers
 - Single-precision floating-point numbers
 - Memory addresses, or pointers
- Unless otherwise stated, we'll assume words are the basic unit of data.



An array of words

- Remember to be careful with memory addresses when accessing words.
- For instance, assume an array of words begins at address 2000.
 - The first array element is at address 2000.
 - The second word is at address 2004, not 2001.
- Revisiting the earlier example, if \$a0 contains 2000, then

```
lw $t0, 0($a0)
```

accesses the first word of the array, but

```
lw $t0, 8($a0)
```

would access the *third* word of the array, at address 2008.

Computing with memory

- So, to compute with memory-based data, you must:
 1. Load the data from memory to the register file.
 2. Do the computation, leaving the result in a register.
 3. Store that value back to memory if needed.
- For example, let's say that you wanted to do the same addition, but the values were in memory. How can we do the following using MIPS assembly language?


```
char A[4] = {1, 2, 3, 4};
```

```
int result;
```

```
result = A[0] + A[1] + A[2] + A[3];
```

Memory alignment

- Keep in mind that memory is byte-addressable, so a 32-bit word actually occupies four contiguous locations (bytes) of main memory.



- The MIPS architecture requires words to be **aligned** in memory; 32-bit words must start at an address that is divisible by 4.
 - 0, 4, 8 and 12 are valid **word addresses**.
 - 1, 2, 3, 5, 6, 7, 9, 10 and 11 are *not* valid word addresses.
 - Unaligned memory accesses result in a **bus error**, which you may have unfortunately seen before.
- This restriction has relatively little effect on high-level languages and compilers, but it makes things easier and faster for the processor.

Next time

- Our next topic is **control flow** in MIPS
 - In section, we'll introduce **loops**
 - Next lecture, we'll introduce **if**, **if/else**, and **case/switch**

How the class will be organized

- **Lecture** and **section** will present course material
 - section useful for gauging your understanding of material
 - weekly, graded on effort, and good practice for the exams
- **Pop quizzes** in-class, based on assigned readings
- **Machine problems** (MPs) more open-ended applications of material
 - can be done in groups (1-3 people)
- **Homeworks** used for closed-form, quantitative problems, before exams
- **Exams** three in-class midterms and one final
- See course info: <http://www.cs.uiuc.edu/class/cs232/html/info.html>

MP 1 is out - due next week!!

- It covers bit-wise operations in C and MIPS (see Section 1 material)
- Go to section - some new material will be presented there!
- Make sure you have a working EWS account soon!!!
- Work in groups (up to 3)

General hints to reach CS232 nirvana

- **Remember the big picture.**
What are we trying to accomplish, and why?
- **Read the lecture notes/textbook.**
Past students have found the lecture notes well organized.
- **Talk to each other.**
You can learn a lot from other CS232 students, both by asking and answering questions. Find some good partners for the homeworks (but make sure you all understand what's going on).
- **Help us help you.**
Lectures, sections, office hours, email, newsgroup. Ask lots of questions!
Check out the web page:

<http://www.cs.uiuc.edu/class/fa06/cs232/>