

Some Real Problem

- What if a program needs more memory than the machine has?
 - even if individual programs fit in memory, how can we run multiple programs?
- How do we protect one program's data from being read or written by another program?
 - multiple programs may want to store something at the same address
 - in particular, consider multiple *copies* of the same program
- There are two key ideas used to solve these problems:
 1. Treat the disk as an extended source of memory
 - swap programs between disk and memory as required
 2. Programs use “fake” or “virtual” memory addresses
 - these translate to “real” addresses, but the translation is hidden to the programmer


Indirection

- Many problems can be solved by adding a level of indirection
- Example: “Which bus will take me from here to IMPE?”
- Without indirection: hard to say, because I don’t know which (physical) buses are running right now
- With indirection: “26 Pack”; this (virtual) name identifies one of several (physical) buses

Without Indirection




With Indirection




Indirection

- Indirection is the ability to reference something using a name or reference instead the value itself.
 - A **mapping** between names and things allows changing the thing without notifying holders of the name.
- **More Examples:**
Pointers, Domain Name Service (DNS) name->IP address, phone system (e.g., cell phone number portability), snail mail (e.g., mail forwarding), 911 (routed to local office), DHCP, color maps, call centers that route calls to available operators, etc.
- In the context of memory:



Virtual Memory


- We translate “virtual addresses” used by the program to “physical addresses” that represent places in the machine’s “physical” memory
 - The word “translate” denotes a level of indirection



A virtual address can be mapped to either physical memory or disk.

Virtual Memory

- Because different processes will have different mappings from virtual to physical addresses, two programs can freely use the same virtual address
- By allocating distinct regions of physical memory to A and B, they are prevented from reading/writing each others data




Caching revisited

- Once the translation infrastructure is in place, the problem boils down to caching.
 - We want the size of disk, but the performance of memory
- The design of virtual memory systems is really motivated by the high cost of accessing disk
 - While memory latency is ~100 times that of cache, disk latency is ~100,000 times that of memory
 - i.e., the miss penalty is HUGE
- Hence, we try to minimize the miss rate:
 - VM “pages” are much larger than cache blocks (why?)
 - A fully associative policy (with approximate LRU) is used
- Should a write-through or write-back policy be used?

Finding the right page

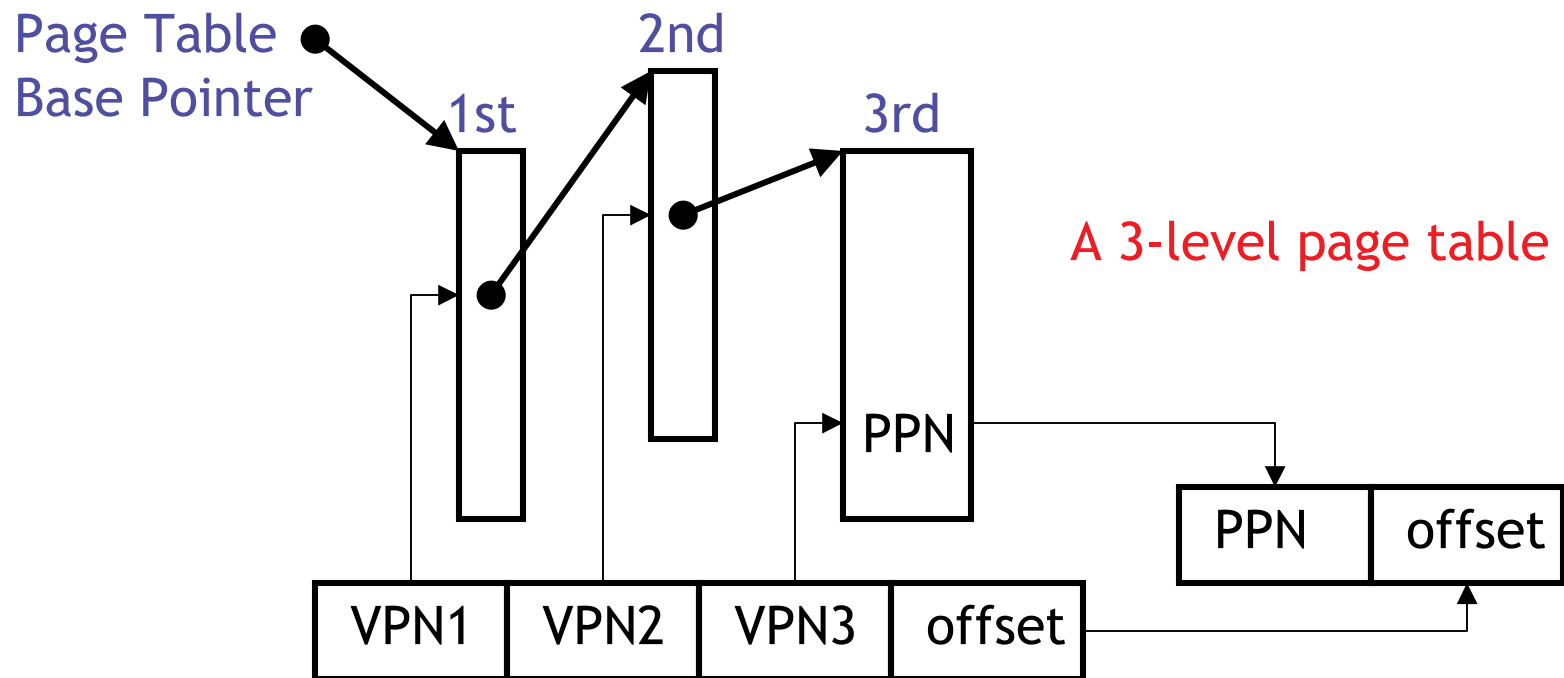
- If it is fully associative, how do we find the right page **without scanning all of memory**?
 - Use an **index** (similar to an index in a book)
- The index is called the **page table**:
 - Each process has a separate page table
 - A “page table register” points to the current process’s page table
 - The page table is indexed with the **virtual page number** (VPN)
 - The VPN is all of the bits that aren’t part of the page offset
 - Each entry contains a valid bit, and a **physical page number** (PPN)
 - The PPN is concatenated with the page offset to get the physical address
 - No tag is needed because the index is the full VPN

Page Table picture



Dealing with large page tables

- Multi-level page tables



- Since most processes don't use the whole address space, you don't allocate the tables that aren't needed
 - Also, the 2nd and 3rd level page tables can be “paged” to disk

Wait a minute!

- We've just replaced every memory access $\text{MEM}[\text{addr}]$ with:

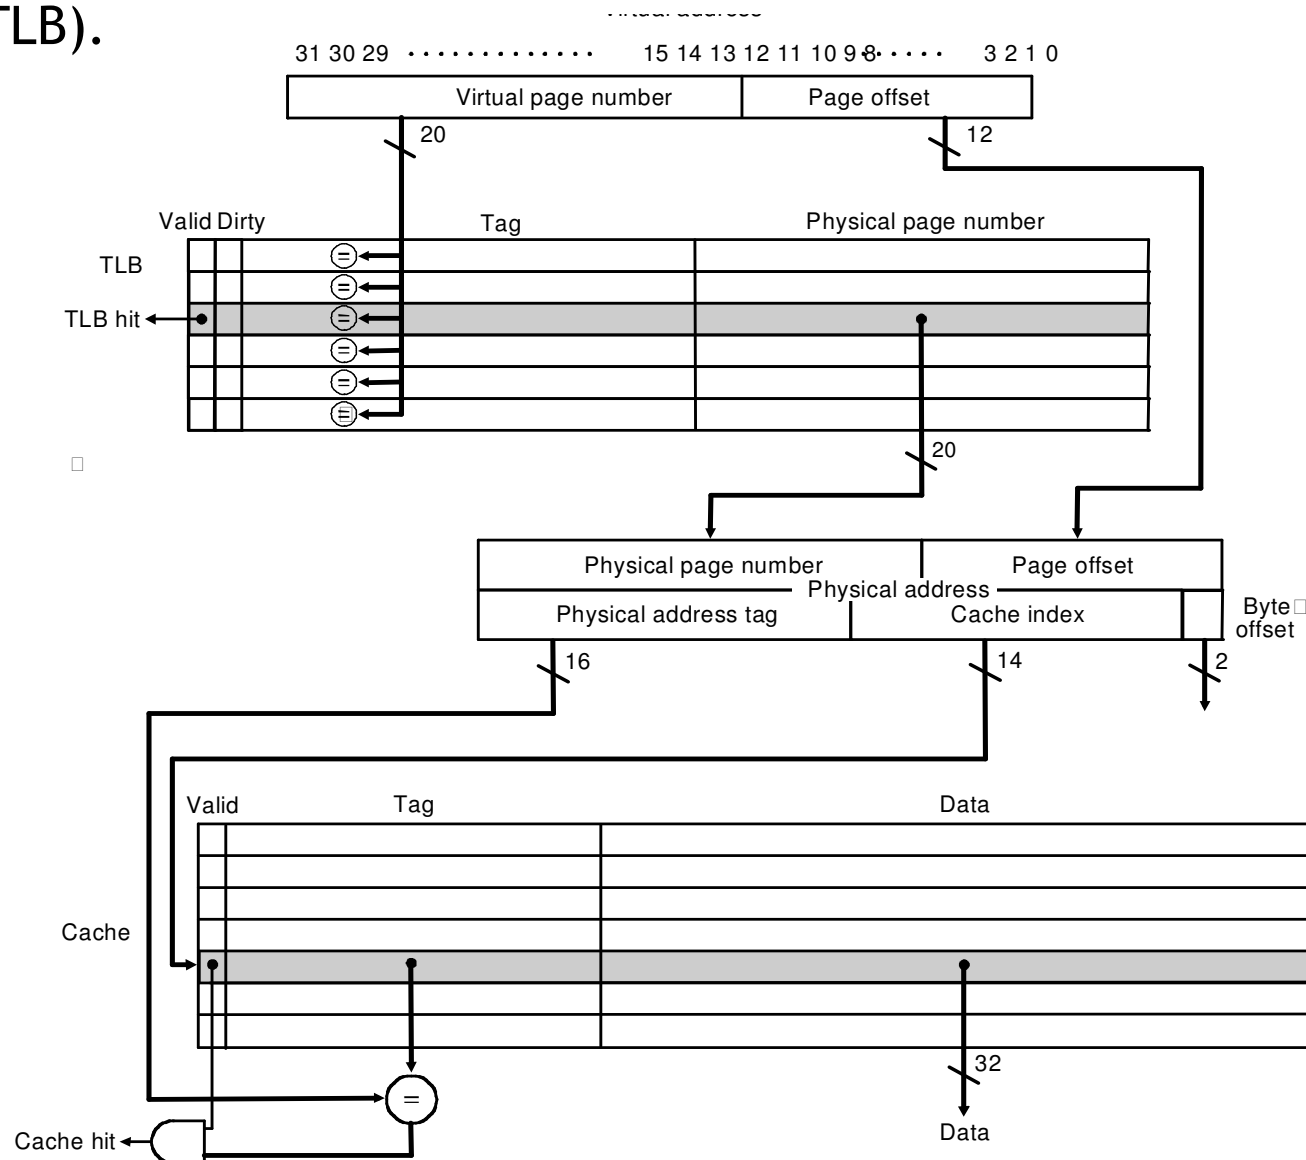
$\text{MEM}[\text{MEM}[\text{MEM}[\text{MEM}[\text{PTBR} + \text{VPN1} \ll 2] + \text{VPN2} \ll 2] + \text{VPN3} \ll 2] + \text{offset}]$

— i.e. 4 memory accesses

- And we haven't talked about the bad case yet (i.e. page faults)...
- We have too many levels of indirection!
- How do we deal with too many levels of indirection?

Caching Translations

- Virtual to Physical translations are cached in a **Translation Lookaside Buffer (TLB)**.



What about a TLB miss?


- If we miss in the TLB, we need to “walk the page table”
 - In MIPS, an exception is raised and software fills the TLB
 - In x86, a “hardware page table walker” fills the TLB
- What if the page is not in memory?
 - This situation is called a **page fault**
 - The operating system will have to request the page from disk
 - It will need to select a page to replace
 - The O/S tries to approximate LRU (see CS423)
 - The replaced page will need to be written back if dirty

Memory Protection

- In order to prevent one process from reading/writing another process's memory, we must ensure that a process cannot change its virtual-to-physical translations
- Typically, this is done by:
 - Having two processor modes: user & kernel
 - Only the O/S runs in kernel mode
 - Only allowing kernel mode to write to the virtual memory state:
 - The page table
 - The page table base pointer
 - The TLB

Sharing Memory

- Paged virtual memory enables sharing at the granularity of a page, by allowing two page tables to point to the same physical addresses
- For example, if you run two copies of a program, the O/S will share the code pages between the programs



Summary

- Virtual memory is **great**:
 - It means that we don't have to manage our own memory
 - It allows different programs to use the same memory
 - It provides protect between different processes
 - It allows controlled sharing between processes (albeit somewhat inflexibly)
- The key technique is **indirection**:
 - Yet another classic CS trick you've seen in this class
 - Many problems can be solved with indirection
- Caching made a few cameo appearances, too:
 - Virtual memory enables using physical memory as a cache for disk
 - We used caching (in the form of the Translation Lookaside Buffer) to make Virtual Memory's indirection fast