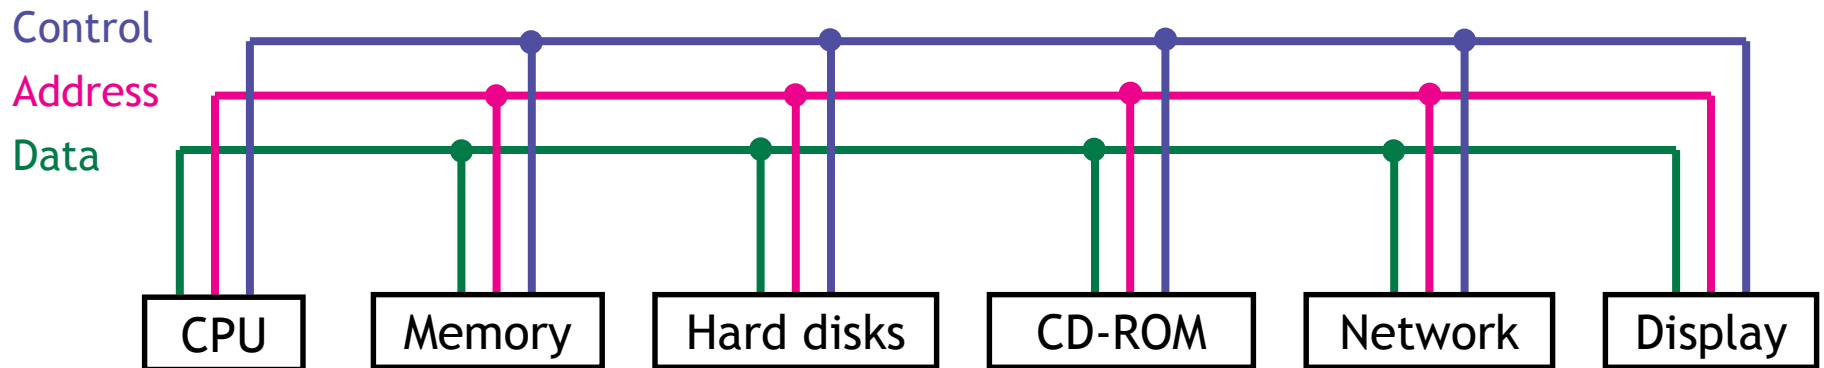


Announcements

- Final exam on Thursday, Dec. 14th from 1:30pm to 4:30pm
 - email me if you need a conflict
- Midterm 3 syllabus:
 - Caches : lectures 15, 16, 17, 18
 - Virtual Memory : lecture 19
 - IO : lectures 21, 22
 - Sections 8 and 9
 - Lecture 20 (parallelism) and Section 10 (VTune) not on Exam 3
- Practice exam will be released middle of next week

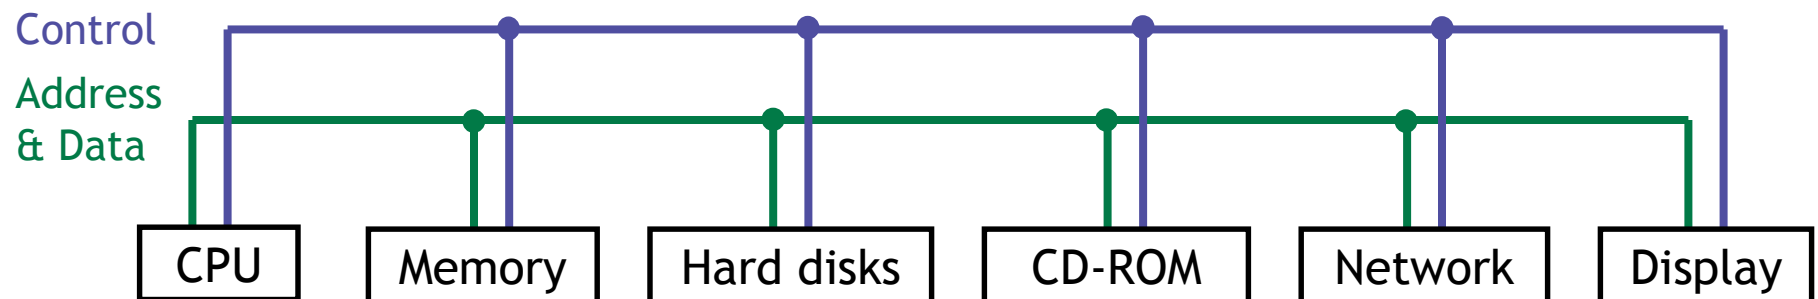
What is the bus anyway?

- A bus is just a bunch of wires which transmits three kinds of information
 - **control signals** (“read” or “write”)
 - the **address** on the device to read or write
 - the actual **data** being transferred
- Some buses have separate control, address and data lines
 - all takes one clock cycle



Multiplexed bus lines

- This is expensive
 - buses transfer 32 to 64 bits of data at a time
 - addresses are usually at least 32-bits long
- Another approach: **multiplex** some lines
 - same lines for address and the data (one after the other)
 - it takes *two* cycles to transmit both pieces of information




Example problem from HW 2, part C

- Interleaved memory, 4 banks (each has latency = 6 cycles)
- Cache block-size = 16-bytes
- CPU-memory bus has latency = 2 cycles and width = 4 bytes

Clean block

Multiplexed address/data bus



Separate address/data buses

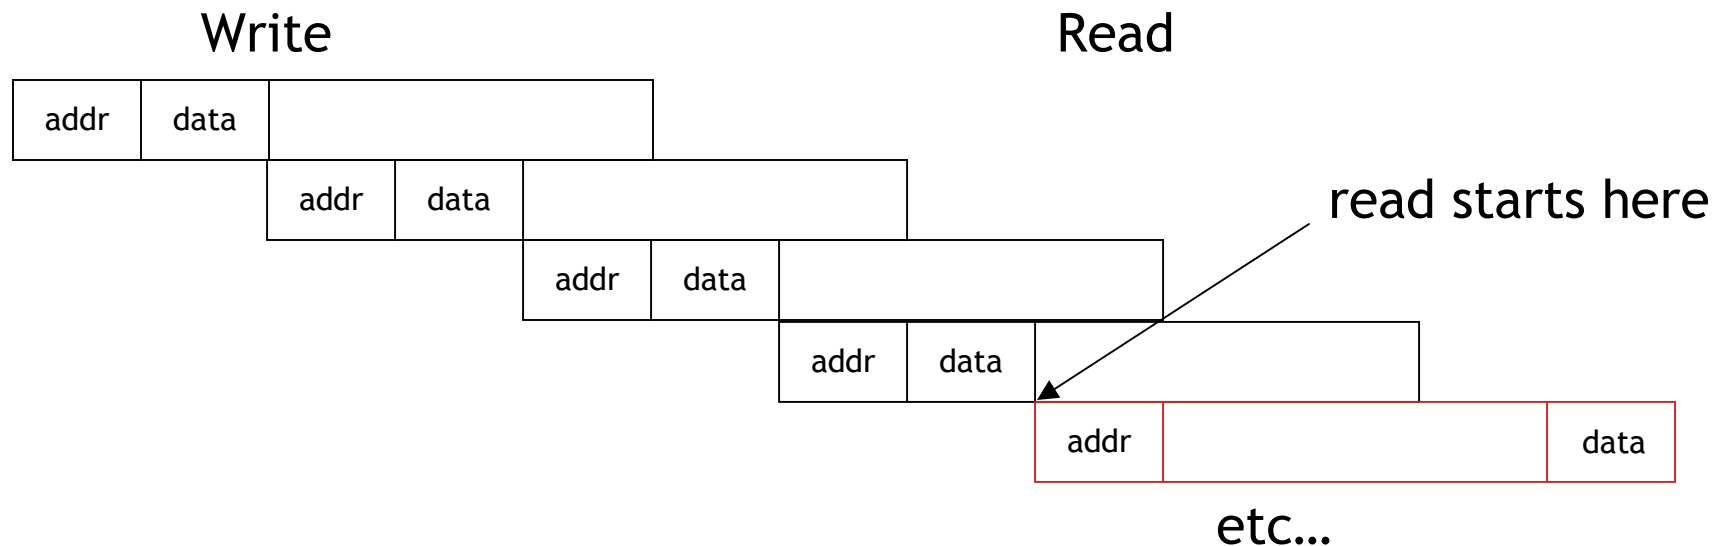
Same picture, but a fifth transfer would not stall since `addr` and `data` use different buses

A fifth transfer would have stalled since `addr` and `data` share the bus

Example problem from HW 2, part C

- Interleaved memory, 4 banks (each has latency = 6 cycles)
- Cache block-size = 16-bytes
- CPU-memory bus has latency = 2 cycles and width = 4 bytes

Dirty block (multiplexed bus)



Frequencies

- CPUs actually operate at two frequencies:
 - The **internal frequency** is the clock rate inside the CPU, which is what we've been talking about so far
 - The **external frequency** is the speed of the processor bus, which limits how fast the CPU can transfer data
- The internal frequency is usually a multiple of the external bus speed
 - A 2.167 GHz Athlon XP sits on a 166 MHz bus (166×13)
 - A 2.66 GHz Pentium 4 might use a 133 MHz bus (133×20)

Synchronous and asynchronous buses

- A **synchronous bus** operates with a central clock signal
 - transactions can be handled easily with finite state machines
 - all devices on the bus run at the same speed (even if they can run faster)
 - limited number of devices (more devices means longer bus)
- An **asynchronous bus** does not rely on clock signals
 - transactions rely on complicated protocols
 - the bus can be longer, devices can operate at different speeds
 - external buses like USB and Firewire are asynchronous

External buses

- **External buses** support the frequent plugging and un-plugging of devices
- **Universal Serial Bus (USB)** and **FireWire**:
 - **Plug-and-play** standards devices configured by software, instead of flipping switches or setting jumpers
 - **Hot plugging** add/remove peripheral without turning machine off
 - The cable transmits **power** - fewer power cables
 - **Serial links** are used, so the cable and connectors are small



The Serial/Parallel conundrum

- Generally, more wires means more bandwidth - quicker transfer times

Recall the equation:

$$\text{Time} = \text{latency} + \text{transfer_size} / \text{bandwidth}$$

- Other things being equal, bigger bandwidth means less time
- But bigger bandwidth (parallel) also means bigger latency...
Serial buses can be clocked at higher frequencies:
 - parallel interconnects have **interference** between signal wires
 - **skew** is also a problem: bits arrive at slightly different times
- Serial links are being increasingly considered for internal buses:
 - **Serial ATA** is a new standard for hard drive interconnects
 - **PCI-Express** (aka 3GI/O) is a PCI bus replacement that uses serial links

Error Correcting Codes

- Physical representations of bits can get corrupted
 - disks can get scratched
 - interference
 - cosmic rays!!
 - ...
- Detect and correct errors by introducing **redundancy**
 - store data plus check-bits
- Example 1: along with every byte, store its **parity**
Instead of storing just 1001 1101, store 1001 1101 **1** ← **parity bit**
If *any* bit is flipped (even the parity bit), error can be detected
- Example 2: store every bit multiple times
Instead of storing just 1, store 1 1 1
Use “majority” to detect and correct upto one error

Redundant bits

- Data range is 0x00, ..., 0xFF (8 bits)
- Instead of storing 8-bit values, we store $8 + r$ bit values
- Can represent 2^{8+r} things, but only 2^8 “real” values
 - codewords (pre-defined list)
- All non-codewords indicate errors
 - correct these to “nearest” codeword
 - Hamming distance between two binary numbers A and B is the number of bits of A to flip until you get B

How much redundancy?

- We want to **detect** upto d errors and **correct** upto c errors
- Codewords must be “spaced out” with a minimum distance of:
 - at least $d + 1$ for “detect” condition
If they were only d apart, one codeword could suffer d errors and be received as another codeword - impossible to detect error!
 - at least $2c + 1$ for “correct” condition
If the closest codeword is at distance c , all other codewords must be at distance at least $c + 1$ in order to determine the *unique* corrected codeword. So, no two codewords are less than $2c + 1$ apart in Hamming distance.
- Take the *maximum* of these two for d bit detection *and* c bit correction
- Common scheme: SECDED (single error correction, double error detection).