


Atomic Operations in Hardware

- Previously, we introduced thread-level parallelism.
 - Today we'll look at instruction support for synchronization.
 - And some pitfalls of parallelization.
 - And solve a few mysteries.



AMD dual-core Opteron

A simple piece of code

```
unsigned counter = 0;
```

```
void *do_stuff(void * arg) {  
    for (int i = 0 ; i < 2000000000 ; ++ i) {  
        counter ++;  
    }  
    return arg;  
}
```

 adds one to counter

How long does this program take? Time for 2000000000 iterations

How can we make it faster? Run iterations in *parallel*

Exploiting a multi-core processor

```
unsigned counter = 0;
```


```
void *do_stuff(void * arg) {  
    for (int i = 0 ; i < 2000000000 ; ++ i) {  
        counter ++;  
    }  
    return arg;  
}
```

Split for-loop across
multiple threads running
on separate cores



NTHREADS

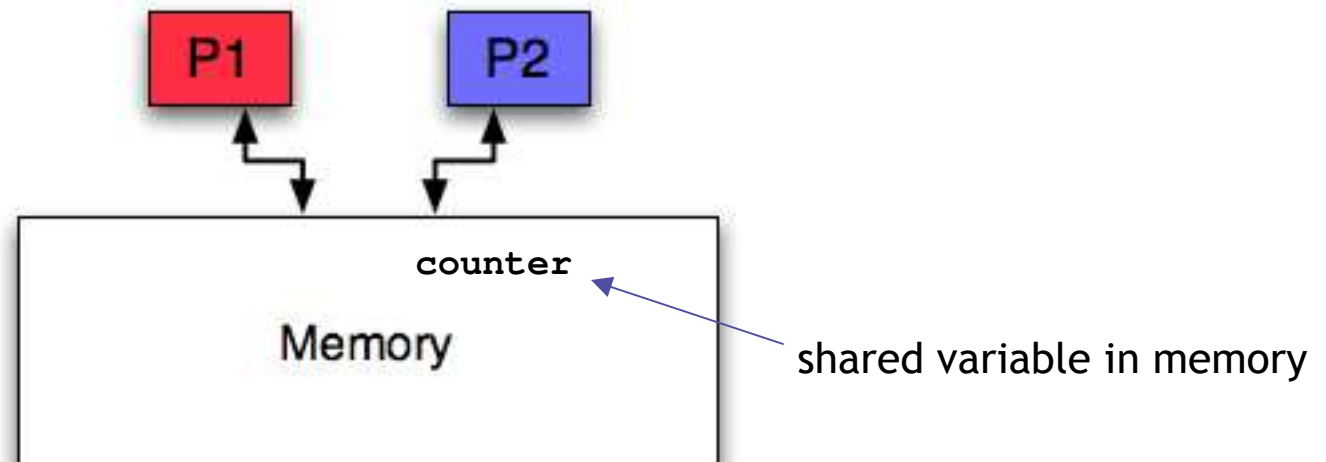
#1



#2

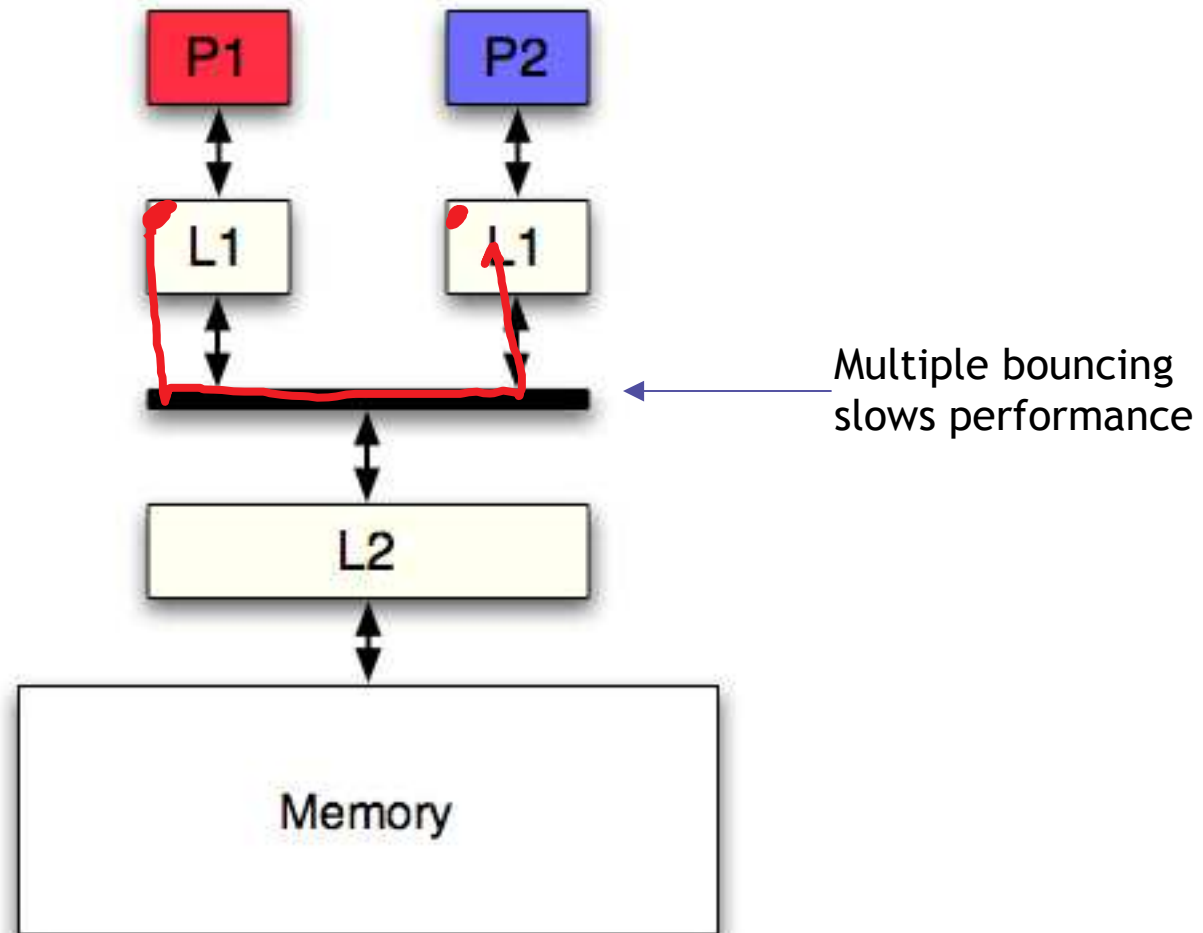
How much faster?

- We're expecting a speedup of 2
- OK, perhaps a little less because of Amdahl's Law
 - overhead for forking and joining multiple threads
- But its actually **slower!!** Why??
- Here's the mental picture that we have - two processors, shared memory



This mental picture is wrong!

- We've forgotten about **caches**!
 - The memory may be shared, but each processor has its own L1 cache
 - As each processor updates `counter`, it bounces between L1 caches



The code is not only slow, its WRONG!

- Since the variable `counter` is *shared*, we can get a **data race**
- Increment operation: `counter++` MIPS equivalent:

```
lw    $t0, counter
addi  $t0, $t0, 1
sw    $t0, counter
```
- A data race occurs when data is **accessed** and **manipulated** by multiple processors, and the outcome depends on the sequence or timing of these events.

Sequence 1

Processor 1

```
lw    $t0, counter
addi  $t0, $t0, 1
sw    $t0, counter
```

Processor 2

```
lw    $t0, counter
addi  $t0, $t0, 1
sw    $t0, counter
```

`counter` increases by 2

Sequence 2

Processor 1

```
lw    $t0, counter
addi  $t0, $t0, 1
sw    $t0, counter
```

Processor 2

```
lw    $t0, counter
addi  $t0, $t0, 1
sw    $t0, counter
```


`counter` increases by 1 !!

Atomic operations

- You can show that if the sequence is particularly nasty, the final value of `counter` may be as little as 2, instead of 200000000.
- To fix this, we must do the load-add-store in a *single* step
 - We call this an **atomic** operation
 - We’re saying: “Do this, and don’t get interrupted while doing this.”
- “Atomic” in this context means “all or nothing”
 - either we succeed in completing the operation with **no interruptions** or we fail to even begin the operation (because someone else was doing an atomic operation)
- x86 provides a “lock” prefix that tells the hardware:
 - “**don’t let anyone read/write the value until I’m done with it**”
 - Not the default case (because it is slow!)

What if we want to generalize beyond increments?

- The lock prefix only works for individual x86 instructions.
- What if we want to execute an arbitrary region of code without interference?
 - Consider a red-black tree used by multiple threads.
- Best mainstream solution: **Locks**
 - Implements **mutual exclusion**
 - You can't have it if I have it, I can't have it if you have it



when lock = 0, set lock = 1, continue

lock = 0

Lock acquire code

- Conceptually, the following code is executed *atomically* to acquire a lock:

High-level version

```
unsigned lock = 0;

while (1) {
    If (lock == 0) {
        lock = 1;
        break;
    }
}
```

MIPS version

```
spin: lw    $t0, 0($a0)
      bne   $t0, 0, spin
      li    $t1, 1
      sw    $t1, 0($a0)
```

- Must be atomic, otherwise we'll get race conditions again!

Race condition in lock-acquire

```
spin: lw    $t0, 0($a0)
      bne   $t0, 0, spin
      li    $t1, 1
      sw    $t1, 0($a0)
```

Lock 10

P0

P1

LW → 10
bne
li
sw 1

LW → 10
bne
li
sw 1

Doing it atomically

- Make sure no one gets between load and store
- Common primitive: **compare-and-swap** (**old**, **new**, **addr**)
 - If the value in memory matches “old”, write “new” into memory

```
temp = *addr;  
if (temp == old) {  
    *addr = new;  
} else {  
    old = temp;  
}
```

← lets us know if we failed

- x86 calls it CMPXCHG (compare-exchange)
 - Use the lock prefix to guarantee it's atomicity

Using CAS to implement locks

- Acquiring the lock:

```
lock_acquire:
```

```
    li    $t0, 0    # old
```

```
    li    $t1, 1    # new
```

```
    cas   $t0, $t1, lock
```

```
    beq   $t0, $t1, lock_acquire    # failed, try again
```

- Releasing the lock:

```
    sw    $t0, lock
```

Conclusions

- When parallel threads access the same data, potential for **data races**
 - Even true on uniprocessors due to context switching
- We can prevent data races by enforcing **mutual exclusion**
 - Allowing only one thread to access the data at a time
 - For the duration of a critical section
- Mutual exclusion can be enforced by locks
 - Programmer allocates a variable to “protect” shared data
 - Program must perform 0→1 transition before data access
 - 1 → 0 transition after
- Locks can be implemented with atomic operations
 - (hardware instructions that enforce mutual exclusion on 1 data item)
 - compare-and-swap
 - If address holds “old”, replace with “new”