


Exploiting a multi-core processor

```
unsigned counter = 0;
```

Split for-loop across
multiple threads running
on separate cores

```
void *do_stuff(void * arg) {  
    parallel_for (int i = 0 ; i < 2000000000 ; ++ i) {  
        counter ++;  
    }  
    return arg;  
}
```

#1



#2

Problem 1: Data Races

- Since `counter` is *shared* between threads, we can get a **data race**

```
parallel_for (int i = 0 ; i < 2000000000 ; ++ i) {  
    counter ++;  
}
```

```
lw    $t0, counter  
addi  $t0, $t0, 1  
sw    $t0, counter
```

Fix problem by making
this operation **atomic**

Sequence 1

Processor 1

```
lw    $t0, counter  
addi  $t0, $t0, 1  
sw    $t0, counter
```

Processor 2

```
lw    $t0, counter  
addi  $t0, $t0, 1  
sw    $t0, counter
```

`counter` increases by 2

Sequence 2

Processor 1

```
lw    $t0, counter  
addi  $t0, $t0, 1  
sw    $t0, counter
```

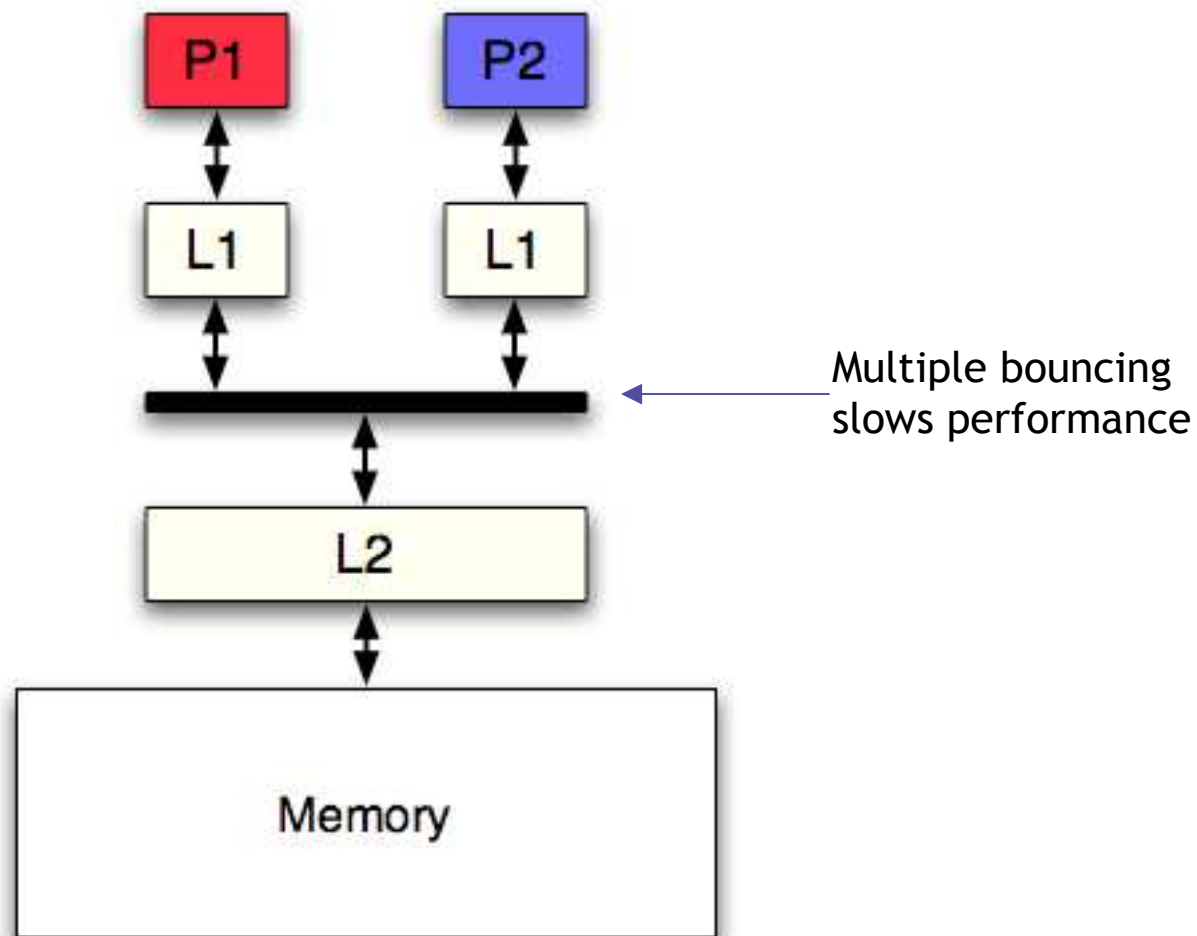
Processor 2

```
lw    $t0, counter  
addi  $t0, $t0, 1  
sw    $t0, counter
```

`counter` increases by 1 !!

Problem 2: This is SLOW!

- Memory is shared, but the L1 caches are not!



Fast, parallel code (software approach)

- Hardware guarantees correctness with **atomic operations**, but its slow

```
parallel_for (i = 0; i < N; ++i) {  
    counter++;  
}
```

- What if each thread had its *own* copy of `counter`? (*private*, not shared)

```
parallel_for (i = 0; i < N; ++i) private(counter) {  
    counter++; // increment local copy  
}  
// Now reduce the local copies of counter into a single variable
```

- This works because “+” is associative and commutative
– fortunately, common operations have these properties

An example: Selection Sort

```
for(i = 0; i < LENGTH - 1; i++) {  
    min = i;          // assume ith element is smallest  
  
    for(j = i + 1; j < LENGTH; j++) {  
        if(a[j] < a[min])  
            min = j;  // oops, jth element is smaller  
    }  
  
    swap(a[min], a[i]);  
}
```

Another example: Register Allocation

- Recall the **register allocation** problem
 - assign register names for every program variable and temporary
 - variables that are **live** at the same time must get different registers
- Graph model
 - Vertices = Variables
 - edge between u and v if both variables are live at the same time
- Register allocation == Graph coloring
 - assign colors to every vertex so that **neighbors** get different colors
- For special graphs, there is a simple optimal coloring algorithm...

Greedy Coloring

```
for all vertices v {
    v.color =  $\infty$ ;
    v.weight = 0;
    bag.insert(v);
}

for(i = 0; i < N; ++i) {
    v = bag.remove_heaviest();
    available_colors = {1, 2, ..., N};

    for all neighbors u of v
        if(u.color ==  $\infty$ )
            bag.update(u, u.weight++);
        else
            available_colors.remove(u.color);

    v.color = available_colors.smallest();
}
```


Greedy Coloring - slightly modified

```
for all vertices v {  
    v.color =  $\infty$ ;  
    v.weight = 0;  
    bag.insert(v);  
}
```

```
for(i = 0; i < N; ++i) {  
    vertex[i] = bag.remove_heaviest();  
    for all neighbors u of vertex[i]  
        if(u.color ==  $\infty$ ) bag.update(u, u.weight++);  
}
```

```
for(i = 0; i < N; ++i) {  
    available_colors = {1, 2, ..., N};  
    for all neighbors u of vertex[i]  
        if(u.color !=  $\infty$ ) available_colors.remove(u.color);  
    vertex[i].color = available_colors.smallest();  
}
```

vertex



Greedy Coloring - slightly modified, rearranged

```
for all vertices v {  
    v.color =  $\infty$ ;  
    v.weight = 0;  
    bag.insert(v);  
}
```

```
for(i = 0; i < N; ++i) {  
    vertex[i] = bag.remove_heaviest();  
    for all neighbors u of vertex[i]  
        if(u.color ==  $\infty$ )  
            bag.update(u, u.weight++);  
}
```


vertex



```
for(i = 0; i < N; ++i) {  
    available_colors = {1, 2, ..., N};  
    for all neighbors u of vertex[i]  
        if(u.color !=  $\infty$ )  
            available_colors.remove(u.color);  
    vertex[i].color =  
        available_colors.smallest();  
}
```

- What does this look like?

A generalization: Actors



- Each actor executes its own code
- Actors communicate via messages
 - *GEN*'s output becomes *COL*'s input
- “Everything is an actor”, rather like “everything is an object”
 - actors act *concurrently*

Conclusions

- The hardware must implement some basic **atomic** operations
 - we can use these to safely share data across processors
- Can't rely too much on atomic operations - they are slow!
 - they force one processor to wait
- Software approaches:
 1. If possible, use **private** (instead of shared) variables in each thread
 - After parallel region, **reduce** private copies into a single variable
 - Common operations are associative + commutative, so this is OK
 2. Pipelining
 - More generally, actors